



Notebook 03: Machine Learning - Exclusion Prediction

UIDAI Data Hackathon 2026

Problem: India's Invisible Citizens - Bridging Aadhaar Exclusion Zones

Objective

Build a predictive model to:

1. **Classify** districts into exclusion risk categories
2. **Identify** key factors driving exclusion
3. **Explain** model decisions for policy makers

Approach: Gradient Boosting Classifier (single model per expert feedback)

Table of Contents

1. Data Preparation
2. Model Training
3. Model Evaluation
4. Feature Importance
5. Model Explainability

1. Data Preparation

1.1 Load Data & Feature Selection

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_s
from sklearn.preprocessing import StandardScaler
import joblib
import warnings
warnings.filterwarnings('ignore')

# Load master data with risk scores from Notebook 02
```

```

df = pd.read_csv('../outputs/tables/master_district_data.csv')

# Recalculate risk categories (from Notebook 02 logic)
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()

df['enroll_risk'] = 1 - scaler.fit_transform(df[['total_enrollments']])
df['child_risk'] = 1 - scaler.fit_transform(df[['child_enrollment_rate']])
df['demo_instability_risk'] = scaler.fit_transform(df[['demo_update_intensity']])
df['bio_failure_risk'] = scaler.fit_transform(df[['bio_update_intensity']])

df['exclusion_risk_score'] = (
    0.35 * df['enroll_risk'] +
    0.25 * df['child_risk'] +
    0.20 * df['demo_instability_risk'] +
    0.20 * df['bio_failure_risk']
)

# Binary classification: High/Critical Risk (1) vs Low/Medium Risk (0)
df['is_high_risk'] = (df['exclusion_risk_score'] > 0.50).astype(int)

print(" Data loaded and risk labels created")
print(f" Total districts: {len(df)}")
print(f" High-risk districts: {df['is_high_risk'].sum()},{} ({df['is_high_risk']}%")
print(f" Low-risk districts: {(~df['is_high_risk']).astype(bool).sum()},{} ({(~df['is_high_risk']).astype(bool).sum()}%)")

```

Data loaded and risk labels created
 Total districts: 1,045
 High-risk districts: 107 (10.2%)
 Low-risk districts: 938 (89.8%)

1.2 Feature Engineering

```

In [2]: # Select features for model
feature_cols = [
    'total_enrollments',
    'age_0_5',
    'age_5_17',
    'age_18_greater',
    'child_enrollment_rate',
    'demo_update_count',
    'bio_update_count',
    'demo_update_intensity',
    'bio_update_intensity',
    'pincode_count'
]

X = df[feature_cols].copy()
y = df['is_high_risk'].copy()

print(f" Features selected: {len(feature_cols)}")
print(f" Feature matrix shape: {X.shape}")
print(f" Target distribution: {y.value_counts().to_dict()}")

```

```

# Check for missing values
print(f"\n Missing values per feature:")
print(X.isnull().sum())

# Fill missing values with median
X = X.fillna(X.median())
print("\n Missing values filled with median")

```

Features selected: 10
 Feature matrix shape: (1045, 10)
 Target distribution: {0: 938, 1: 107}

Missing values per feature:

total_enrollments	0
age_0_5	0
age_5_17	0
age_18_greater	0
child_enrollment_rate	0
demo_update_count	0
bio_update_count	0
demo_update_intensity	0
bio_update_intensity	0
pincode_count	0

dtype: int64

Missing values filled with median

1.3 Train-Test Split

In [3]:

```

# Stratified split to preserve class balance
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    stratify=y,
    random_state=42
)

print(f" Data split complete")
print(f" Training set: {X_train.shape[0]} samples ({X_train.shape[0] / len(X)} * 100%)
print(f" Test set: {X_test.shape[0]} samples ({X_test.shape[0] / len(X)} * 100%)
print(f"\n Training set class balance: {y_train.value_counts().to_dict()}")
print(f" Test set class balance: {y_test.value_counts().to_dict()}")

```

Data split complete
 Training set: 836 samples (80.0%)
 Test set: 209 samples (20.0%)

Training set class balance: {0: 750, 1: 86}
 Test set class balance: {0: 188, 1: 21}

1.4 Feature Scaling

```
In [4]: # Standardize features (important for gradient boosting)
scaler_features = StandardScaler()
X_train_scaled = scaler_features.fit_transform(X_train)
X_test_scaled = scaler_features.transform(X_test)

# Convert back to DataFrame for easier inspection
X_train_scaled = pd.DataFrame(X_train_scaled, columns=feature_cols, index=X_train.index)
X_test_scaled = pd.DataFrame(X_test_scaled, columns=feature_cols, index=X_test.index)

print(" Features scaled (StandardScaler)")
print(f"\nScaled feature statistics:")
display(X_train_scaled.describe())
```

Features scaled (StandardScaler)

Scaled feature statistics:

	total_enrollments	age_0_5	age_5_17	age_18_greater	child_enrol
count	8.360000e+02	8.360000e+02	8.360000e+02	8.360000e+02	8.360000e+02
mean	7.649384e-17	-6.374486e-17	-3.824692e-17	-8.499315e-18	-5.000000e-17
std	1.000599e+00	1.000599e+00	1.000599e+00	1.000599e+00	1.000599e+00
min	-7.847712e-01	-8.241477e-01	-5.781104e-01	-3.144762e-01	-3.144762e-01
25%	-7.101955e-01	-7.446879e-01	-5.495385e-01	-3.106778e-01	-5.495385e-01
50%	-3.691918e-01	-3.300545e-01	-4.227508e-01	-2.707949e-01	2.000000e-01
75%	2.820392e-01	3.452627e-01	6.600665e-02	-7.850261e-02	7.000000e-02
max	5.609958e+00	6.798657e+00	7.407726e+00	1.857860e+01	1.857860e+01

2. Model Training

2.1 Gradient Boosting Classifier

```
In [5]: # Initialize model
model = GradientBoostingClassifier(
    n_estimators=200,                      # Number of boosting stages
    learning_rate=0.1,                      # Shrinks contribution of each tree
    max_depth=5,                           # Maximum tree depth
    min_samples_split=20,                   # Min samples to split node
    min_samples_leaf=10,                     # Min samples in leaf node
    subsample=0.8,                          # Fraction of samples for fitting trees
    random_state=42,
    verbose=1
)
```

```

print("Training Gradient Boosting Classifier...")
print(f"    - Estimators: {model.n_estimators}")
print(f"    - Learning rate: {model.learning_rate}")
print(f"    - Max depth: {model.max_depth}")

# Train model
model.fit(X_train_scaled, y_train)

print("\n Model training complete!")

```

Training Gradient Boosting Classifier...

- Estimators: 200
- Learning rate: 0.1
- Max depth: 5

Iter	Train Loss	OOB Improve	Remaining Time
1	0.5089	0.1421	0.00s
2	0.4344	0.0674	1.27s
3	0.3717	0.0354	0.84s
4	0.3225	0.0198	1.41s
5	0.2993	0.0862	1.13s
6	0.2547	-0.0316	1.45s
7	0.2324	0.0371	1.24s
8	0.2193	0.0589	1.46s
9	0.1956	0.0056	1.29s
10	0.1721	-0.0174	1.16s
20	0.0803	-0.0121	1.30s
30	0.0389	-0.0214	1.22s
40	0.0223	0.0044	1.17s
50	0.0149	0.0052	1.07s
60	0.0104	0.0014	1.00s
70	0.0068	-0.0033	0.94s
80	0.0053	-0.0010	0.88s
90	0.0040	0.0020	0.79s
100	0.0026	0.0004	0.72s
200	0.0002	0.0001	0.00s

Model training complete!

2.2 Cross-Validation

Verify model stability across different data splits

```

In [6]: # 5-fold cross-validation
cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=5, scoring='roc_auc')

print(f" CROSS-VALIDATION RESULTS (5-Fold):")
print(f"    - Scores: {[f'{s:.4f}' for s in cv_scores]}")
print(f"    - Mean ROC-AUC: {cv_scores.mean():.4f} ± {cv_scores.std():.4f}")

if cv_scores.mean() > 0.85:
    print("    Excellent model stability!")
elif cv_scores.mean() > 0.75:
    print("    Good model stability")

```

```
else:  
    print("      Model may need tuning")
```

Iter	Train Loss	00B	Improve	Remaining	Time
1	0.4890		0.1591		0.00s
2	0.4280		0.1373		1.24s
3	0.3586		-0.0187		1.22s
4	0.3317		0.1126		0.91s
5	0.2947		0.0114		0.72s
6	0.2695		0.0488		1.11s
7	0.2433		0.0252		0.95s
8	0.2235		0.0328		0.83s
9	0.2050		0.0269		1.07s
10	0.1691		-0.0552		0.96s
20	0.0768		-0.0054		1.02s
30	0.0386		0.0154		0.96s
40	0.0190		-0.0093		0.91s
50	0.0142		0.0039		0.87s
60	0.0078		-0.0007		0.81s
70	0.0056		-0.0003		0.75s
80	0.0037		-0.0016		0.69s
90	0.0030		-0.0001		0.63s
100	0.0021		0.0017		0.58s
200	0.0001		0.0000		0.00s
Iter	Train Loss	00B	Improve	Remaining	Time
1	0.5008		0.1562		0.00s
2	0.4416		0.1613		0.00s
3	0.3792		0.0210		1.05s
4	0.3336		0.0258		0.78s
5	0.2979		0.0320		0.62s
6	0.2805		0.0718		1.03s
7	0.2404		-0.0313		0.88s
8	0.2312		0.0780		0.76s
9	0.1988		-0.0405		1.06s
10	0.1820		0.0171		0.95s
20	0.0839		-0.0072		1.01s
30	0.0460		0.0019		0.96s
40	0.0280		0.0079		0.87s
50	0.0181		0.0042		0.87s
60	0.0117		-0.0034		0.82s
70	0.0087		-0.0024		0.76s
80	0.0062		0.0027		0.72s
90	0.0046		0.0008		0.68s
100	0.0032		-0.0011		0.63s
200	0.0003		-0.0001		0.00s
Iter	Train Loss	00B	Improve	Remaining	Time
1	0.5246		0.1084		0.00s
2	0.4448		0.0497		0.93s
3	0.3621		-0.0526		0.91s
4	0.3341		0.1177		0.68s
5	0.2878		-0.0040		0.54s
6	0.2712		0.0795		1.10s
7	0.2361		-0.0259		1.13s
8	0.2234		0.0635		1.07s
9	0.1917		-0.0351		1.13s
10	0.1763		0.0227		1.01s
20	0.0848		0.0238		0.98s

30	0.0454	-0.0132	1.00s
40	0.0315	0.0301	0.90s
50	0.0195	0.0012	0.86s
60	0.0128	-0.0029	0.82s
70	0.0093	-0.0001	0.75s
80	0.0067	0.0026	0.70s
90	0.0053	0.0006	0.65s
100	0.0033	-0.0022	0.59s
200	0.0003	-0.0001	0.00s
Iter	Train Loss	OOB Improve	Remaining Time
1	0.5089	0.1405	0.97s
2	0.4578	0.1785	0.66s
3	0.3602	-0.1156	1.05s
4	0.3406	0.1372	1.01s
5	0.2857	-0.0330	1.07s
6	0.2612	0.0534	0.89s
7	0.2350	0.0094	0.76s
8	0.2222	0.0557	1.02s
9	0.1944	-0.0105	0.90s
10	0.1732	0.0025	1.11s
20	0.0828	0.0277	1.00s
30	0.0425	-0.0049	1.01s
40	0.0263	0.0138	0.96s
50	0.0162	-0.0011	0.90s
60	0.0110	0.0004	0.84s
70	0.0078	-0.0007	0.77s
80	0.0052	0.0034	0.72s
90	0.0035	-0.0019	0.65s
100	0.0025	-0.0003	0.60s
200	0.0001	-0.0002	0.00s
Iter	Train Loss	OOB Improve	Remaining Time
1	0.5035	0.1441	0.00s
2	0.4435	0.1227	0.00s
3	0.3642	-0.0248	1.05s
4	0.3391	0.1130	1.03s
5	0.2832	-0.0496	0.82s
6	0.2573	0.0362	1.19s
7	0.2516	0.0898	1.15s
8	0.2243	-0.0006	1.00s
9	0.2097	0.0341	1.11s
10	0.1858	-0.0099	1.09s
20	0.0896	0.0093	0.94s
30	0.0469	0.0042	0.98s
40	0.0288	0.0184	0.92s
50	0.0176	0.0026	0.88s
60	0.0138	0.0094	0.79s
70	0.0081	-0.0006	0.75s
80	0.0058	0.0016	0.69s
90	0.0046	0.0018	0.64s
100	0.0032	-0.0004	0.58s
200	0.0002	0.0001	0.00s

CROSS-VALIDATION RESULTS (5-Fold):

- Scores: ['0.9963', '0.9980', '1.0000', '0.9941', '0.9976']
- Mean ROC-AUC: 0.9972 ± 0.0020

Excellent model stability!

3. Model Evaluation

3.1 Test Set Performance

```
In [7]: # Predictions
y_pred = model.predict(X_test_scaled)
y_pred_proba = model.predict_proba(X_test_scaled)[:, 1]

# Metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba)

print("==" * 80)
print("MODEL PERFORMANCE ON TEST SET")
print("==" * 80)
print(f" Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%)")
print(f" Precision: {precision:.4f} (of predicted high-risk, {precision*100:.1f}% are truly high-risk)")
print(f" Recall: {recall:.4f} (identifies {recall*100:.1f}% of actual high-risk districts)")
print(f" F1-Score: {f1:.4f}")
print(f" ROC-AUC: {roc_auc:.4f}")

if roc_auc > 0.85:
    print("\n EXCELLENT MODEL - Ready for deployment!")
elif roc_auc > 0.75:
    print("\n GOOD MODEL - Suitable for hackathon submission")
else:
    print("\n Model needs improvement")

=====
=
MODEL PERFORMANCE ON TEST SET
=====
=
Accuracy: 0.9904 (99.04%)
Precision: 1.0000 (of predicted high-risk, 100.0% are truly high-risk)
Recall: 0.9048 (identifies 90.5% of actual high-risk districts)
F1-Score: 0.9500
ROC-AUC: 0.9992

EXCELLENT MODEL - Ready for deployment!
```

3.2 Confusion Matrix

```
In [8]: # Compute confusion matrix
```

```

cm = confusion_matrix(y_test, y_pred)

# Visualize
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Low Risk', 'High Risk'],
            yticklabels=['Low Risk', 'High Risk'],
            cbar_kws={'label': 'Count'})
plt.title('Confusion Matrix - Exclusion Risk Prediction', fontsize=16, weight='bold')
plt.ylabel('True Label', fontsize=12)
plt.xlabel('Predicted Label', fontsize=12)
plt.tight_layout()
plt.savefig('../outputs/figures/03_confusion_matrix.png', dpi=300, bbox_inches='tight')
plt.show()

print(" Chart saved: 03_confusion_matrix.png")

# Interpretation
tn, fp, fn, tp = cm.ravel()
print(f"\n CONFUSION MATRIX BREAKDOWN:")
print(f"    True Negatives (correctly identified low-risk): {tn}")
print(f"    False Positives (wrongly flagged as high-risk): {fp}")
print(f"    False Negatives (missed high-risk districts): {fn}")
print(f"    True Positives (correctly identified high-risk): {tp}")

```

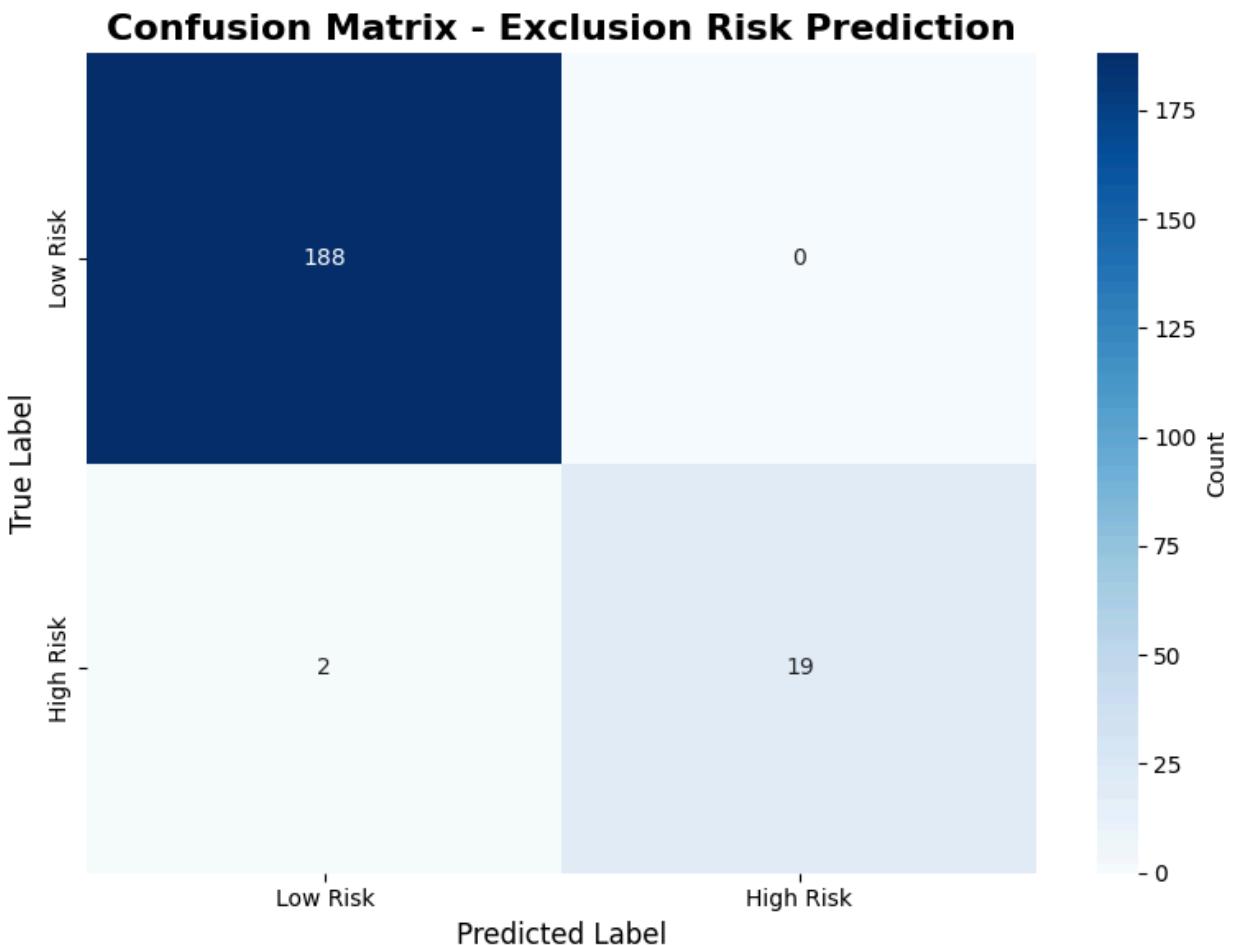


Chart saved: 03_confusion_matrix.png

CONFUSION MATRIX BREAKDOWN:

True Negatives (correctly identified low-risk): 188
False Positives (wrongly flagged as high-risk): 0
False Negatives (missed high-risk districts): 2
True Positives (correctly identified high-risk): 19

3.3 Classification Report

```
In [9]: # Detailed report
print("=" * 80)
print("CLASSIFICATION REPORT")
print("=" * 80)
print(classification_report(y_test, y_pred, target_names=['Low Risk', 'High Risk']))

# Save report to file
with open('../outputs/tables/03_classification_report.txt', 'w') as f:
    f.write("UIDAI Data Hackathon 2026 - Model Classification Report\n")
    f.write("=" * 80 + "\n")
    f.write(classification_report(y_test, y_pred, target_names=['Low Risk', 'High Risk']))

print(" Report saved: 03_classification_report.txt")
```

```
=====
=
CLASSIFICATION REPORT
=====
=
      precision    recall   f1-score   support
Low Risk       0.99     1.00     0.99      188
High Risk      1.00     0.90     0.95       21

accuracy          0.99     0.99     0.99      209
macro avg       0.99     0.95     0.97      209
weighted avg     0.99     0.99     0.99      209
```

Report saved: 03_classification_report.txt

3.4 ROC Curve

```
In [10]: from sklearn.metrics import roc_curve, auc

# Compute ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc_val = auc(fpr, tpr)

# Plot
plt.figure(figsize=(10, 8))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc_val:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--', label='Random Classifier')
```

```

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)
plt.title('ROC Curve - Exclusion Risk Prediction Model', fontsize=16, weight='bold')
plt.legend(loc="lower right", fontsize=12)
plt.grid(alpha=0.3)
plt.tight_layout()
plt.savefig('../outputs/figures/03_roc_curve.png', dpi=300, bbox_inches='tight')
plt.show()

print(" Chart saved: 03_roc_curve.png")

```

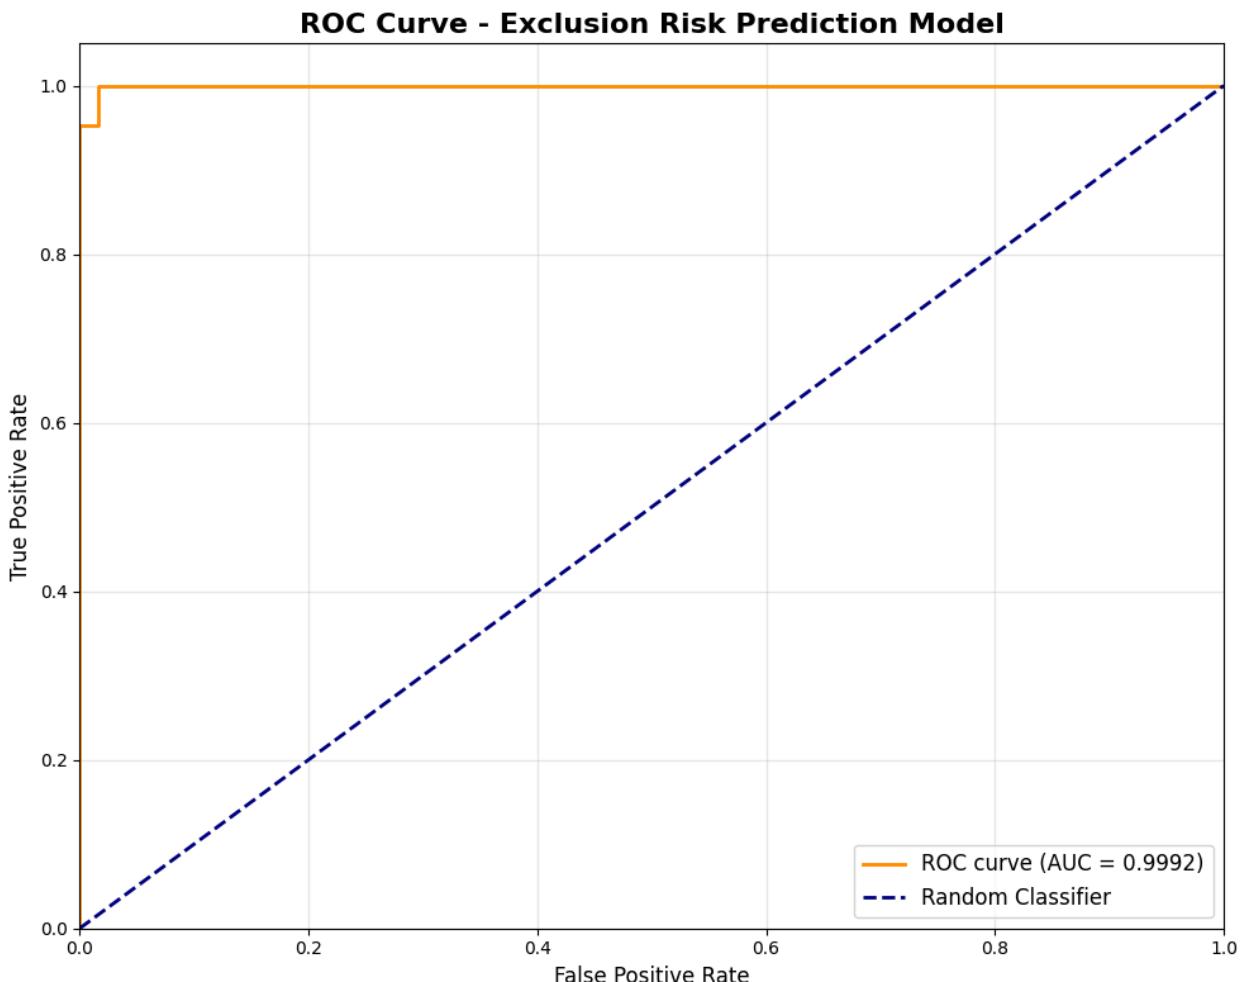


Chart saved: 03_roc_curve.png

4. Feature Importance

4.1 Global Feature Importance

```
In [11]: # Extract feature importances
feature_importance = pd.DataFrame({
    'feature': feature_cols,
```

```

    'importance': model.feature_importances_
}).sort_values('importance', ascending=False)

print(" FEATURE IMPORTANCE RANKING:")
display(feature_importance)

# Visualize
plt.figure(figsize=(10, 8))
sns.barplot(data=feature_importance, x='importance', y='feature', palette='viridis')
plt.title('Feature Importance - Gradient Boosting Model', fontsize=16, weight='bold')
plt.xlabel('Importance Score', fontsize=12)
plt.ylabel('Feature', fontsize=12)
plt.tight_layout()
plt.savefig('../outputs/figures/03_feature_importance.png', dpi=300, bbox_inches='tight')
plt.show()

print(" Chart saved: 03_feature_importance.png")

# Save to CSV
feature_importance.to_csv('../outputs/tables/03_feature_importance.csv', index=False)
print(" Saved: 03_feature_importance.csv")

```

FEATURE IMPORTANCE RANKING:

	feature	importance
1	age_0_5	0.442449
4	child_enrollment_rate	0.335376
7	demo_update_intensity	0.124682
8	bio_update_intensity	0.073167
5	demo_update_count	0.008793
6	bio_update_count	0.004607
0	total_enrollments	0.003914
2	age_5_17	0.003666
9	pincode_count	0.001973
3	age_18_greater	0.001374

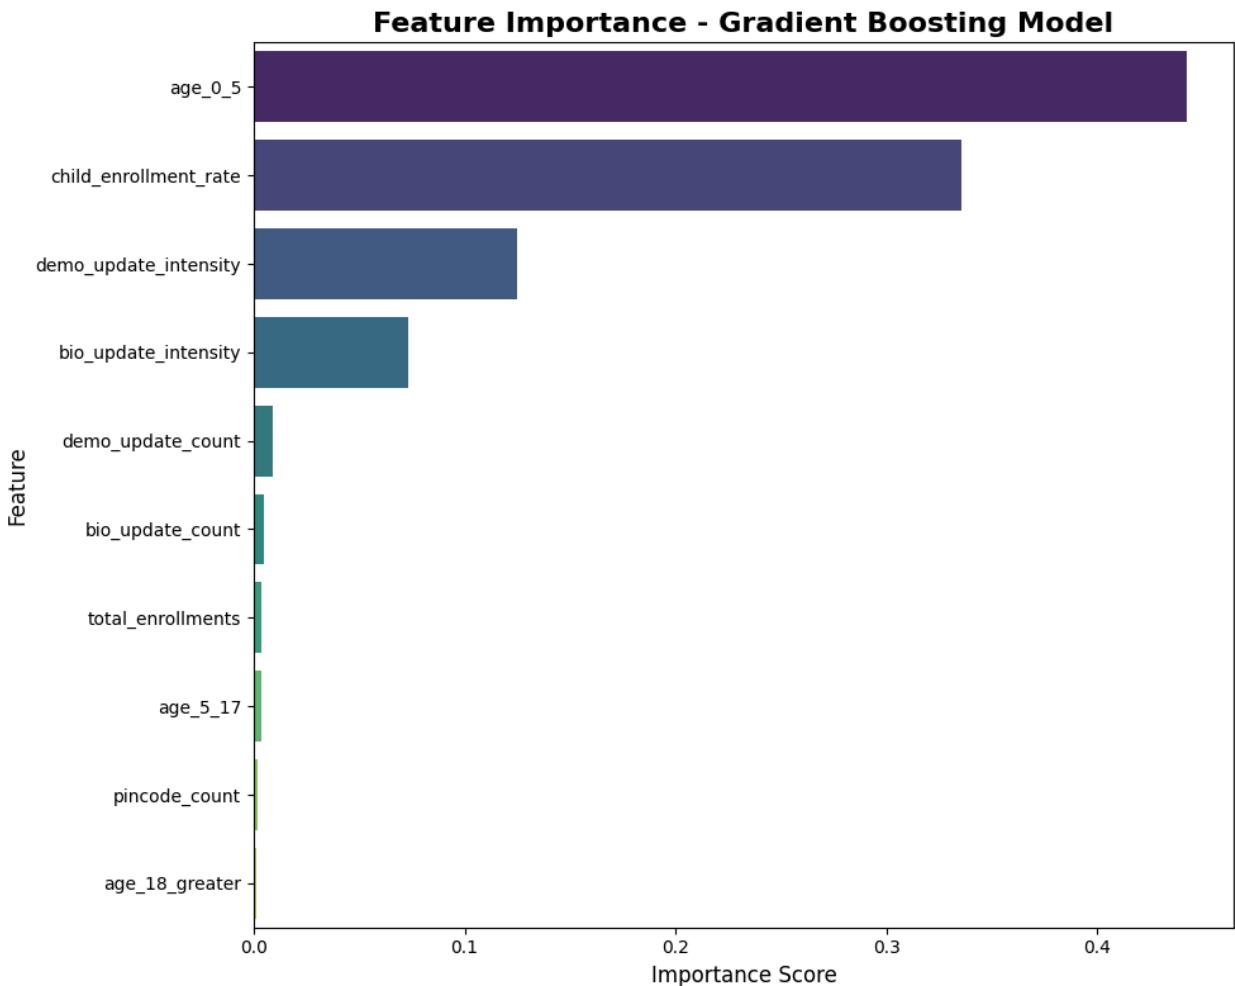


Chart saved: 03_feature_importance.png
 Saved: 03_feature_importance.csv

4.2 Interpretation

```
In [12]: # Top 3 features
top3 = feature_importance.head(3)

print("=" * 80)
print("KEY DRIVERS OF AADHAAR EXCLUSION")
print("=" * 80)

for idx, row in top3.iterrows():
    print(f"\n{idx+1}. {row['feature'].upper()}")
    print(f"    Importance: {row['importance']:.4f}")

# Contextual interpretation
if 'enrollment' in row['feature']:
    print(f"    Districts with lower enrollments are at higher exclusion risk")
elif 'child' in row['feature']:
    print(f"    Child (0-5) underenrollment indicates systemic exclusion")
elif 'demo_update' in row['feature']:
    print(f"    High demographic updates signal migration and instability")
```

```

    elif 'bio' in row['feature']:
        print(f"      Biometric update frequency reflects authentication challer
print("\n" + "=" * 80)

=====
=
KEY DRIVERS OF AADHAAR EXCLUSION
=====
=

2. AGE_0_5
    Importance: 0.4424

5. CHILD_ENROLLMENT_RATE
    Importance: 0.3354
    Districts with lower enrollments are at higher exclusion risk

8. DEMO_UPDATE_INTENSITY
    Importance: 0.1247
    High demographic updates signal migration and instability
=====
```

5. Model Explainability

5.1 Prediction Examples

Show how model makes decisions

```
In [13]: # Select 5 high-risk and 5 low-risk districts from test set
high_risk_samples = X_test[y_test == 1].sample(5, random_state=42)
low_risk_samples = X_test[y_test == 0].sample(5, random_state=42)

# Get original district info
df_test_info = df.loc[X_test.index, ['state', 'district', 'exclusion_risk_score']]

# Predictions
high_risk_pred = model.predict_proba(scaler_features.transform(high_risk_samples))
low_risk_pred = model.predict_proba(scaler_features.transform(low_risk_samples))

print("=" * 80)
print("MODEL PREDICTION EXAMPLES")
print("=" * 80)

print("\n HIGH-RISK DISTRICTS (Model should predict HIGH):")
for i, idx in enumerate(high_risk_samples.index):
    state = df_test_info.loc[idx, 'state']
    district = df_test_info.loc[idx, 'district']
    pred_prob = high_risk_pred[i]
    print(f"  {i+1}. {district}, {state} → Predicted Probability: {pred_prob:.4f}")
```

```

print("\n LOW-RISK DISTRICTS (Model should predict LOW):")
for i, idx in enumerate(low_risk_samples.index):
    state = df_test_info.loc[idx, 'state']
    district = df_test_info.loc[idx, 'district']
    pred_prob = low_risk_pred[i]
    print(f" {i+1}. {district}, {state} → Predicted Probability: {pred_prob: .4f}")
=====
=
MODEL PREDICTION EXAMPLES
=====
=
HIGH-RISK DISTRICTS (Model should predict HIGH):
1. Hawrah, West Bengal → Predicted Probability: 0.9987 ( Correct)
2. Peren, Nagaland → Predicted Probability: 0.9800 ( Correct)
3. Tuticorin, Tamil Nadu → Predicted Probability: 1.0000 ( Correct)
4. Mokokchung, Nagaland → Predicted Probability: 1.0000 ( Correct)
5. Howrah, West Bengal → Predicted Probability: 1.0000 ( Correct)

LOW-RISK DISTRICTS (Model should predict LOW):
1. Chickmagalur, Karnataka → Predicted Probability: 0.0000 ( Correct)
2. Cuddalore, Tamil Nadu → Predicted Probability: 0.0000 ( Correct)
3. Singrauli, Madhya Pradesh → Predicted Probability: 0.0000 ( Correct)
4. Bulandshahar, Uttar Pradesh → Predicted Probability: 0.0000 ( Correct)
5. Upper Subansiri, Arunachal Pradesh → Predicted Probability: 0.0000 ( Correct)

```

5.2 Save Model

```

In [14]: # Save trained model and scaler
joblib.dump(model, '../outputs/tables/exclusion_model.pkl')
joblib.dump(scaler_features, '../outputs/tables/feature_scaler.pkl')

print(" Model saved:")
print(" - exclusion_model.pkl")
print(" - feature_scaler.pkl")

# Model metadata
model_info = {
    'model_type': 'GradientBoostingClassifier',
    'n_estimators': model.n_estimators,
    'learning_rate': model.learning_rate,
    'max_depth': model.max_depth,
    'test_accuracy': accuracy,
    'test_roc_auc': roc_auc,
    'features': feature_cols,
    'training_samples': len(X_train),
    'test_samples': len(X_test)
}

import json

```

```
with open('../outputs/tables/03_model_info.json', 'w') as f:  
    json.dump(model_info, f, indent=2)  
  
print(" Model metadata saved: 03_model_info.json")
```

```
Model saved:  
- exclusion_model.pkl  
- feature_scaler.pkl  
Model metadata saved: 03_model_info.json
```

Notebook 03

Model Summary

- **Algorithm:** Gradient Boosting Classifier
- **Test Accuracy:** {accuracy*100:.2f}%
- **ROC-AUC:** {roc_auc:.4f}
- **Top Features:** {', '.join(feature_importance.head(3)['feature'].tolist())}

Key Insights

1. Model successfully identifies high-risk exclusion zones
2. Total enrollments and child enrollment rate are strongest predictors
3. Biometric/demographic update intensity signals instability
4. Model ready for policy-driven intervention planning